

Job Aggregation Platform

Final Project Report

Prepared for: Engineering & Product Management

Prepared by: Perna Anand

Date: July 6, 2026

Dataset analyzed: jobs.json — 155,794 unique job records

Source documents: research/api_analysis.md, reports/analytics_summary.md

Table of Contents

Cover Page
2. Executive Summary
3. Project Overview
4. Dataset Overview
5. Provider Analysis
6. Regional Coverage
7. Job Freshness
8. Duplicate Analysis
9. API Analysis
10. India Research
11. Additional Regional APIs
12. API Filter Optimization
13. Maximum Coverage from Existing Providers
14. Engineering Recommendations
15. Roadmap
16. Future Improvements
17. Conclusion

2. Executive Summary

The Job Aggregation Platform aggregates 155,794 unique job postings from 20 registered data providers into a single normalized schema, deduplicated by URL, and served through a static frontend. Nineteen of the 20 providers are currently contributing data; one (TheirStack) is blocked by an account plan restriction unrelated to the codebase.

Coverage is broad but concentrated. The United States accounts for 51% of all jobs, and a single provider — SmartRecruiters — supplies 41.8% of the entire dataset. A single employer, Domino's, alone accounts for 15.7% of every job in the platform, because its board lists each store-level opening as a separate posting across thousands of locations.

Deduplication is functioning correctly: zero duplicate URLs remain in the dataset. A large volume of same-title, same-company postings initially appears duplicate-like, but inspection confirms these are legitimate multi-location roles (retail, logistics, delivery) rather than aggregation errors — a naive dedup rule would have destroyed tens of thousands of real job listings.

Freshness is moderate but degrading. At collection time, 54.0% of jobs were posted within the prior 30 days; 46.0% are already stale, and 14.0% are over a year old. The pipeline has no scheduled refresh and no expiry mechanism, so this ratio will continue to worsen without intervention.

The highest-value opportunities identified are: (1) extending the existing, already-keyed Adzuna integration to additional countries — unlocking Canada, the UK, continental Europe, and Latin America with no new integration work; (2) introducing a scheduled refresh and expiry cycle to fix freshness; (3) closing the Singapore/Southeast Asia and India coverage gaps via a no-auth government API and existing ATS infrastructure; and (4) turning on date and location filters already supported by several integrated APIs but currently unused. None of these require scraping, new commercial contracts, or code changes to the aggregation pipeline's core logic.

3. Project Overview

The platform is a modular Python pipeline (`jobs.py` plus a `sources/` package) that fetches job postings from every registered provider, normalizes each record into a common seven-field schema, deduplicates the combined list by URL, and writes the result to `jobs.json`. A lightweight static frontend (HTML/CSS/JavaScript) reads that file directly to render a searchable, filterable job board.

Schema

Every job record — regardless of source — is normalized to exactly seven fields: title, company, location, url, tags, remote, and posted (an ISO YYYY-MM-DD date string). This uniform shape is what allows 20 structurally different provider APIs to feed a single downstream pipeline and frontend.

Provider architecture

Providers were added in three phases: Phase 1 (6 no-authentication public APIs), Phase 2 (8 API-key-based providers, credentials loaded from a git-ignored `.env` file), and Phase 3 (6 company-based ATS integrations — Greenhouse, Lever, Ashby, SmartRecruiters, Workable, and Teamtailor — each fetching one job board per individually verified company). Every provider implements the same two-method interface (`fetch_raw`, `normalize`), so adding a new source never requires touching the core pipeline.

Reliability

Each provider's `run()` method isolates failures: a network error, bad response, or per-record normalization issue in one source is logged and skipped rather than aborting the entire run. Retryable HTTP statuses (429, 500, 502, 503, 504) are retried with exponential backoff before a request is given up on.

4. Dataset Overview

The current dataset (jobs.json) contains 155,794 unique job records spanning 9,497 distinct companies. Schema completeness is very high: only 45 records (0.03%) are missing a company name, only 72 (0.05%) are missing a location, and 155,779 records (99.99%) carry a parseable posting date.

Metric	Value
Total unique jobs	155,794
Registered providers	20
Providers contributing in latest run	19
Distinct companies	9,497
Jobs with empty company	45 (0.03%)
Jobs with empty location	72 (0.05%)
Jobs with empty tags	11,762 (7.5%)
Jobs with a parseable posted date	155,779 (99.99%)
Duplicate URLs remaining	0

Remote vs. onsite

Only 10.9% of jobs (16,992) are marked remote; the remaining 89.1% (138,802) are onsite or not marked remote. Several providers are remote-only by construction (Himalayas, RemoteOK, Jobicy); most others derive the remote flag from a genuine API field where available, falling back to a keyword heuristic on title/location otherwise.




Figure 1. Remote vs. onsite split across the full dataset.

Company distribution

Company concentration is extreme and bimodal. The top 10 companies account for 37.8% of the dataset (58,922 jobs); the top 25 account for 47.5%. At the same time, 58.0% of all 9,497 distinct companies (5,505 companies) contribute exactly one job each. Domino's alone represents 15.7% of every job in the platform — nearly one in six listings — because its SmartRecruiters board lists each store-level opening (delivery driver, assistant manager, etc.) as a separate posting across thousands of U.S. locations.

Top Hiring Companies

Top 15 of 9,497 distinct companies · n = 155,794




Figure 2. Top 15 hiring companies by job count.

5. Provider Analysis

Provider attribution was reconstructed from each job's URL host, since the schema does not carry an explicit source field. A single provider, SmartRecruiters, supplies 41.8% of the entire dataset — more than four in ten jobs. The top three providers (SmartRecruiters, Lever, Greenhouse) together account for 62.8% of all jobs.

Provider	Jobs	% of dataset
SmartRecruiters	65,193	41.8%
Lever	17,755	11.4%
Greenhouse	14,930	9.6%
USAJOBS	10,000	6.4%
Reed	9,000	5.8%
Bundesagentur (+AMS)	~7,800	~5.0%
Ashby	5,171	3.3%
Workable	4,159	2.7%
Jooble	~3,100	~2.0%
Himalayas	2,964	1.9%
Adzuna	2,001	1.3%
The Muse	2,000	1.3%
Arbeitnow	930	0.6%
Teamtailor	474	0.3%
Other (SerpApi, OpenWeb Ninja, Careerjet, RemoteOK, Jobicy)	~350	~0.2%
TheirStack	0	0.0% (plan-gated, HTTP 402)




Figure 3. Jobs by provider (top 12, remainder grouped as Other).

ATS provider mix

The six company-based ATS providers (Greenhouse, Lever, Ashby, SmartRecruiters, Workable, Teamtailor) together supply 107,664 jobs — 69.1% of the entire dataset. Within that ATS-sourced subset, SmartRecruiters alone represents 60.6%, followed by Lever (16.5%) and Greenhouse (13.9%); Ashby, Workable, and Teamtailor together make up the remaining 9.1%.




Figure 4. Distribution across the six ATS providers only.

All six ATS providers grow exclusively by manually adding verified company slugs to a curated list — there is no query-based way to expand their reach. Regional and company composition is therefore driven largely

by which companies were curated into these lists, not by a balanced global crawl.

6. Regional Coverage

Region was inferred from each job's free-text location field using a keyword/token classifier. The United States dominates at 51.0% of the dataset; the six primary target regions together account for 81.3% (126,638 jobs).

Region	Jobs	% of dataset
United States	79,427	51.0%
Europe (continental, non-UK)	21,886	14.0%
United Kingdom	10,305	6.6%
LATAM	6,955	4.5%
Singapore / Southeast Asia	4,629	3.0%
Canada	3,436	2.2%
Remote / Worldwide	822	0.5%
Other / Unclassified	28,334	18.2%




Figure 5. Jobs by region.

Canada is the weakest of the six primary target regions at 2.2% — under half the size of the UK's dedicated coverage. 18.2% of jobs carry location text too ambiguous to classify into a region, which limits the reliability of any location-based filter, count, or map feature in its current form.

Top countries

Rank	Country	Jobs	%
1	United States	79,427	51.0%
2	Germany	12,167	7.8%
3	United Kingdom	10,305	6.6%
4	India	3,471	2.2%
5	Canada	3,436	2.2%
6	France	2,218	1.4%
7	China	1,435	0.9%
8	Australia	1,225	0.8%




Figure 6. Top 15 resolved countries by job count.

Europe is largely a Germany story: Germany's 12,167 jobs represent 56% of the entire European total, driven by the two German public-agency sources. India, at 2.2% (3,471 jobs), is already the fourth-largest country in the dataset despite having no dedicated source — see Section 10.

7. Job Freshness

Age is measured against the dataset's collection date (2026-07-03, the newest posted value present in the data). At collection time, a slim majority of the dataset was current, but the stale fraction is already substantial and will only grow without a refresh mechanism.

Age at collection	Jobs	Cumulative 'within'	Cumulative %
Within 24 hours	24,052	24,052	15.4%
Within 3 days	+6,582	30,634	19.7%
Within 7 days	+11,884	42,518	27.3%
Within 14 days	+24,610	67,128	43.1%
Within 30 days	+16,949	84,077	54.0%
Older than 30 days	71,702	—	46.0%
Unknown / unparseable	15	—	0.0%

Job Freshness

Age measured against collection date 2026-07-03 · n = 155,794




Figure 7. Job freshness distribution (buckets from freshest to stalest).

Total current jobs (≤ 30 days): 84,077 — 54.0% of the dataset. Total stale jobs (> 30 days, including unknown): 71,717 — 46.0%.

Deep-stale tail

Age	Jobs	%
Older than 90 days	45,075	28.9%
Older than 180 days	29,640	19.0%
Older than 365 days	21,768	14.0%

The oldest posting on record dates to 2009-12-05; the median posting date is 2026-06-09. The stale tail is concentrated in large enterprise ATS boards, which return every currently open requisition, including long-standing evergreen roles. Because the dataset is a static snapshot with no scheduled refresh and no expiry logic, the current/stale ratio degrades every day after collection with no offsetting mechanism.

8. Duplicate Analysis

The pipeline deduplicates by exact URL, and this is functioning correctly: zero duplicate URLs and zero empty URLs remain in the dataset. Under looser candidate keys, however, a large number of same-title/company rows appear to “collide” — this section quantifies that finding and explains why it is not, in fact, a duplication problem.

Check	Result
Exact-URL duplicates remaining	0
Jobs with empty URL	0
Title + company collisions (extra rows)	35,300 across 11,340 groups
Title + company + location collisions (extra rows)	8,290 across 4,963 groups
Title+company groups spanning more than one distinct URL	11,340 (100% of those groups)

Duplicate Analysis

*Not true duplicates — same-title/company rows at different URLs, overwhelmingly legitimate multi-location postings (see report).




Figure 8. Duplicate counts under three candidate keys.

Critically, every one of the 11,340 title+company groups spans multiple distinct URLs — these are not missed duplicates. Inspection shows they are overwhelmingly the same role posted at hundreds of separate physical locations: Domino's “Delivery Driver,” BoxLunch “Sales Associate,” and similar multi-location retail/logistics roles. Applying a naive title+company deduplication rule would have erased tens of thousands of legitimate, distinct job openings.

The real, and much smaller, duplicate risk is cross-provider overlap between aggregator-style sources (e.g., Jooble, Careerjet, and Adzuna occasionally re-listing the same third-party posting) — a risk the current URL-keyed dedup does not directly measure and any future duplicate-reduction effort should target conservatively, without collapsing on title+company alone.

9. API Analysis

All 20 registered providers were audited for authentication model, rate and result-window limits, pagination behavior, and available filters. Three sources are already capped by the provider's own hard result window (USAJOBS and Bundesagentur at 10,000 jobs; Reed near its ~9,900-job ceiling). Three more are deliberately throttled to control cost on metered/paid APIs (SerpApi, OpenWeb Ninja, and TheirStack, each self-capped around 50 jobs per run).

Priority	Providers	Rationale
Critical	SmartRecruiters, Adzuna	SmartRecruiters is already the volume backbone; Adzuna is the single highest-leverage untapped asset (one keyed integration can add several regions by iterating its country parameter).
High	Lever, Greenhouse, USAJOBS, Reed, Bundesagentur	Large, free/near-free backbones for major regions, though each is bounded by a curated company list, federal-only scope, or a hard result window.
Medium	Ashby, Workable, Himalayas, The Muse, Jooble, Careerjet	Solid contributors with real breadth but individually smaller or with untapped upside (e.g., Himalayas' richer search endpoint).
Low	Arbeitsnow, RemoteOK, Jobicy, Teamtailor, SerpApi, OpenWeb Ninja, TheirStack	Small absolute volume or deliberately throttled metered sources; TheirStack currently contributes zero jobs.

Full per-provider detail — authentication, pricing, rate limits, pagination, and every supported filter — is documented in [research/api_analysis.md](#), Sections 4 and 9.

10. India Research

India is already the 4th-largest country in the dataset (3,471 jobs, 2.2%) despite having no dedicated source. A comprehensive survey of the Indian recruitment ecosystem — government portals, private job boards, ATS/HR platforms, staffing firms, and regional/niche boards — found a consistent structural pattern: India has a large, active online hiring market but almost no aggregatable public job APIs.

Key findings

- Government: the National Career Service (NCS) publishes vacancy data via data.gov.in as a periodic open dataset, not a live searchable API — authoritative but freshness-limited.
- Private boards: Naukri, Foundit, Apna, Shine, TimesJobs, Instahyre, Hirist, iimjobs, Internshala, Unstop, Wellfound, and LinkedIn/Indeed India all expose no official public read API.
- India-origin ATS/HRMS platforms (Zoho Recruit, Freshteam, Darwinbox, Keka, GreythHR) are uniformly per-tenant — useful only to that specific employer's own integration, not aggregatable across employers.
- Major Indian staffing firms (TeamLease, Quess, Randstad India, etc.) expose no public job feeds.

Recommended routes (ranked)

Rank	Route	Effort	ROI
1	ATS slug expansion for Indian-HQ companies on Greenhouse/Lever/Ashby/Workable/SmartRecruiters (already integrated)	Low	High
2	India-scoped query on existing metered sources (JSearch/OpenWeb Ninja, SerpApi)	Low	High
3	Paid aggregator (Fantastic Jobs / Active Jobs DB) covering Indian ATS and boards	Low	Medium (budget-gated)
4	National Career Service open dataset (data.gov.in)	Medium (ETL)	Medium
5	Enterprise data providers (Coresignal, JobsPikr)	Medium	Low (cost)
—	Scraping Naukri / Foundit / Apna / LinkedIn / Indeed	—	Not recommended (legal/maintenance risk)

Bottom line: legal, low-maintenance India coverage comes from extending infrastructure already built for this project (ATS slugs and existing metered/Google-Jobs sources), not from integrating a new Indian ATS vendor or scraping consumer boards.

11. Additional Regional APIs

A verified market survey was conducted across Canada, the UK, Europe, Singapore, Southeast Asia, and LATAM to identify additional free or public data sources beyond the 20 already integrated.

Region	Best candidate	Effort	Priority
Canada	Adzuna 'ca' (existing key, config-level change)	Trivial	Critical
United Kingdom	Adzuna 'gb' (de-risks single-source dependency on Reed)	Trivial	High
Europe	Adzuna multi-country (fr, de, es, nl, it, at, pl)	Trivial	Critical
Europe	France Travail (official, free, OAuth2 registration)	Medium	High
Singapore	MyCareersFuture — official, no-auth public API (~80k jobs)	Low	Critical
Southeast Asia	ATS slug expansion for SEA-headquartered companies	Low	Medium
LATAM	Adzuna 'br' + 'mx' (existing key)	Trivial	Critical
LATAM	Get on Board — the one LATAM board with a real public API	Low–Medium	Investigate

A separate verification pass explicitly ruled out 20 well-known regional boards — Eluta, Workopolis, Jobillico, Talent.com, WowJobs (Canada); TotalJobs, CV-Library, Jobsite, CWJobs (UK); StepStone, Xing Jobs, Jobrapido, Welcome to the Jungle (Europe); Computrabajo, Bumeran, OCC Mundial (LATAM); and JobStreet, JobsDB, Kalibrr, Glints (Southeast Asia) — all of which are either closed with no public API, or offer only a paid partner/posting channel rather than a free read feed. Get on Board (LATAM) was the sole exception found with a genuine public API.

The headline finding: the Adzuna key the project already holds is currently querying only the United States. Extending it to additional countries is, by a wide margin, the single highest-leverage change available — it touches five of the six target regions without any new integration work.

12. API Filter Optimization

An audit of all 20 integrations found that most run deliberately unfiltered “browse everything” queries — the right choice for a first pass, but one that leaves three significant levers unused.

- Date filters (e.g., Adzuna max_days_old, USAJOBS DatePosted, Jooble datecreatedfrom) are supported by at least six integrated providers but used by none — turning them on would cut the stale tail at ingest rather than after the fact.
- Country/locale filters (Adzuna's country path, Careerjet's locale_code, Himalayas' country parameter) are supported but mostly unused — turning them on directly targets the weak regions identified in Section 6.
- Query partitioning — splitting one broad query into several filtered ones by category, location, or date range — is the only way to exceed the hard result-window caps already reached by USAJOBS (10,000), Bundesagentur (10,000), Reed (~9,900), The Muse (2,000), Jooble (~1,100), and Himalayas (3,000 of a 90,000-job archive). This is the single largest coverage lever available without adding any new provider.

Rough planning estimate: turning on these unused filters and partitioning the capped sources could add approximately 40,000 to 90,000 additional useful, fresher, and less-duplicated jobs without integrating a single new provider. Full per-provider filter detail is in research/api_analysis.md, Section 17.

13. Maximum Coverage from Existing Providers

Every one of the 20 registered providers was audited individually to answer a narrower question than the rest of this report: is each provider we already have wired up already returning everything it realistically can, given its own API's own rules — and if not, how much more is obtainable using only capabilities that provider's API already exposes?

Are we already extracting the maximum realistic number of jobs from each provider?

For 17 of the 20 providers, no — meaningful additional volume is available using capabilities already present in the provider's own API, without any new authentication, scraping, or commercial agreement. Only three providers (Arbeitnow, RemoteOK, and TheirStack) are already at their practical ceiling, and one of those three (TheirStack) is capped by an account billing restriction rather than by the API itself.

Which providers still have untapped potential?

- Adzuna — by far the largest opportunity in the entire provider set. The current 2,000-job figure reflects querying a single country (the US) out of 10+ countries the existing API key already has access to.
- Bundesagentur, USAJOBS, and Reed — each already sits at or near a hard per-query result-window ceiling (10,000, 10,000, and ~9,900 respectively), but each also supports a location or keyword filter that can open an independent window of similar size when the query is partitioned instead of run once, broadly.
- Himalayas, The Muse, Jooble, and Careerjet — each is currently run as a single broad or wildcard query well below its real ceiling; switching to filtered, partitioned queries (by country, category, or real keywords instead of a wildcard) would meaningfully increase volume.

Which providers have already reached their practical limits?

- Arbeitnow — the entire current board is already fetched to exhaustion via full pagination, and the API offers no keyword, location, or category parameter to partition further.
- RemoteOK — the documented public endpoint is a fixed ~100-job snapshot with no confirmed filtering or pagination mechanism.
- TheirStack — effectively at zero due to an account plan restriction (HTTP 402), not because its API has been exhausted; this should be reassessed the moment Jobs API access is restored on the account.

Which providers can be improved using existing API filters?

Eight providers — Adzuna, Bundesagentur, USAJOBS, Reed, Himalayas, Jooble, Careerjet, and The Muse — already support date, country/locale, category, or location filters that the current integration does not use. Turning these on, or partitioning a single broad query into several filtered ones, is a pure configuration and query-logic change with no new credentials required.

Which providers require adding additional companies rather than changing code?

All six ATS providers — Greenhouse, Lever, Ashby, SmartRecruiters, Workable, and Teamtailor — expose no keyword, location, category, or date request filter at all; each is a full-board, per-company fetch. For these six, the API itself is already fully exploited for every currently configured company, so the entire opportunity lies in expanding the curated company-slug list, not in changing how the API is queried.

Which providers cannot realistically return more jobs due to API limitations?

Arbeitsnow and RemoteOK. Both expose a complete, unfilterable snapshot of their respective boards with no query dimension to expand beyond what is already being collected — any further volume would depend on the boards themselves growing over time, not on anything this project can change.

Provider-by-provider comparison

Provider	Current Jobs Collected	Estimated Practical Maximum	Additional Jobs Possible	Recommended Optimization Strategy
Arbeitsnow	930	~950–1,000	~0–100	None — full board already exhausted
Himalayas	3,000	13,000–23,000	+10,000–20,000	Switch to richer search endpoint; partition by country/category
RemoteOK	100	~100	~0	None confirmed; verify tag-scoped endpoints before assuming more
Jobicity	100	300–500	+200–400	Partition by geo/industry/tag across multiple calls
The Muse	2,000	5,000–8,000	+3,000–6,000	Partition by category/location; consider an API key
Bundesagentur	10,000	25,000–40,000	+15,000–30,000	Partition by region or keyword
Jobble	1,129	4,000–9,000	+3,000–8,000	Replace wildcard with targeted keywords/locations
USAJOBS	10,000	20,000–30,000	+10,000–20,000	Partition by job category or location
Adzuna	2,000	22,000–42,000	+20,000–40,000	Query additional countries already covered by the existing key
Reed	9,000	17,000–24,000	+8,000–15,000	Partition by UK region/keyword
SerpApi	50	~2,500/mo (budget-bound)	Budget-dependent	Reallocate quota to targeted location/date queries
OpenWeb Ninja	~47	~200/mo (budget-bound)	Budget-dependent	Reallocate quota toward country-scoped (e.g. India) queries
Careerjet	2,000	7,000–12,000	+5,000–10,000	Add locale-scoped passes + real keywords
TheirStack	0	0 (plan-gated)	0 until plan restored	Resolve the account/billing restriction first
Greenhouse	22,728	27,728–37,728	+5,000–15,000	Expand curated company-slug list
Lever	17,755	22,755–27,755	+5,000–10,000	Expand curated company-slug list
Ashby	5,171	7,171–10,171	+2,000–5,000	Source additional slugs outside AI/SaaS/FinTech
SmartRecruiters	65,193	70,193–80,193	+5,000–15,000	Expand slug list, but weigh against concentration risk
Workable	7,644	10,644–14,644	+3,000–7,000	Source slugs in staffing/consulting/fintech/healthcare
Teamtailor	474	974–1,974	+500–1,500	Source additional career sites (highest candidate hit rate)

Summary

- Highest ROI providers to optimize first: Adzuna (+20,000–40,000, no new integration, just additional countries on the existing key), followed by Bundesagentur, USAJOBS, Reed, Himalayas, Jooble, and Careerjet — together representing roughly 56,000–99,000 further jobs from filters and query partitioning these APIs already support.
- Providers already operating near their maximum: Arbeitnow and RemoteOK have no remaining query-based lever; TheirStack is at zero solely due to a billing restriction rather than an API ceiling.
- Expected increase in coverage before integrating additional APIs: on the order of 75,000–140,000 additional jobs are realistically obtainable from the eight already-integrated, already-credentialed providers with untapped filter capacity — before any new provider from the regional or India research in this report needs to be added.

14. Engineering Recommendations

- Broaden regional balance — particularly Canada, UK resilience, and Singapore/Southeast Asia — by extending already-integrated providers rather than adding new integrations.
- Introduce a scheduled refresh and expiry mechanism so the freshness ratio stops degrading between runs and closed postings are retired automatically.
- Grow ATS company-slug lists for under-represented regions (Canada, SEA, LATAM, India) to rebalance provider and regional concentration.
- Do not deduplicate on title+company alone — doing so would remove tens of thousands of legitimate multi-location jobs. Any future duplicate-reduction work should combine URL normalization with conservative, location-aware matching.
- Improve location-text normalization to shrink the 18.2% unclassifiable bucket and make regional filtering and reporting more reliable.
- Revisit the TheirStack account/plan to restore its contribution, or formally retire it from the active provider list if the restriction is permanent.
- Monitor concentration risk — with 41.8% of jobs from one provider and 15.7% from one company, any change to SmartRecruiters' Domino's board would materially shift dataset-wide statistics.
- Turn on already-supported date and country/locale filters (Section 12) before considering any new integration — this is free, low-risk, and high-impact.

15. Roadmap

Recommendations are grouped into five priorities, each with an estimated coverage increase, engineering effort, and business value.

Priority	Focus	Net-new jobs	Effort	Business value
P1	Adzuna multi-country + date filter	15,000–25,000	Low	Very High
P2	Freshness foundation (scheduling + expiry)	0 (quality fix)	Medium	Very High
P3	Weak regions: MyCareersFuture + ATS slugs + India query	40,000–70,000	Low–Medium	High
P4	Result-window bypass via partitioning + France Travail	20,000–50,000	Medium	High
P5	Dedup/location normalization + optional paid aggregator	Quality (+50,000 optional)	Medium	Medium–High

If only two weeks of engineering time were available

- Days 1–2: Ship Adzuna multi-country + max_days_old (Priority 1).
- Days 2–3: Capture a fetched_at timestamp and surface a “data as of” signal.
- Days 3–6: Build the scheduling + presence-diff expiry core (Priority 2).
- Days 6–9: Add MyCareersFuture (Singapore) and a first batch of Indian/SEA/Canada ATS slugs (Priority 3).
- Days 9–12: Prove out one result-window partition (e.g., Reed or The Muse) plus URL normalization in dedup.
- Days 12–14: Integration testing, dedup verification, and rate-limit/budget tuning across the newly added passes.

This sequence realistically adds ~40,000–70,000 net-new, better-distributed jobs and fixes freshness within two weeks, without any scraping or new commercial contracts.

16. Future Improvements

Beyond the two-week roadmap, several structural investments would compound the platform's value as the dataset grows:

- A persistent datastore (e.g., Postgres or OpenSearch) replacing the single jobs.json file, paired with a paginated, queryable search API — the current design of loading the entire file into the browser will not scale past the low hundreds of thousands of jobs.
- Incremental/delta ingestion with per-job first_seen/last_seen tracking and automated expiry, replacing the current full-overwrite snapshot model.
- Scalable deduplication — normalized-URL keys plus blocked fuzzy/locality-sensitive-hash matching — to handle cross-provider overlap at higher volume, and description-based matching once job descriptions are added to the schema.
- Automated ATS company/slug discovery and health monitoring, so coverage grows without linear manual curation effort.
- A decoupled, per-source pipeline with queueing and rate-limit budgeting, so one slow or large provider cannot stall an entire run, with observability on volume, freshness, and duplicate rate per source per run.
- Optional adoption of a paid aggregator (e.g., Fantastic Jobs' Active Jobs DB) for a step-change in both volume and freshness once budget allows.

17. Conclusion

The Job Aggregation Platform is a mature, well-engineered first version: a clean, extensible source architecture spanning 20 providers behind one shared interface, a fully normalized schema, correct URL-based deduplication, and disciplined per-source error isolation. It is production-ready as a periodic bulk job board, but not yet as a freshness- or scale-sensitive product — it remains a static snapshot with no scheduling, no expiry, and no persistent per-job state.

Regional strength is concentrated in the United States, Germany, and (for its size) the UK; Canada and Southeast Asia are thin, India has no dedicated source despite meaningful incidental volume, and the UK's real coverage rides on a single provider. Duplicate handling is excellent at the URL level, and the apparent duplication in title+company groupings is explained, not a defect. Freshness is the platform's weakest dimension and its highest-value fix, being primarily an architectural gap rather than a sourcing one.

Nearly all of the highest-impact opportunities identified in this report are already unlocked: multi-country Adzuna, MyCareersFuture, France Travail, ATS slug expansion, and India coverage via existing infrastructure are all free or low-cost, legally clean, and require no new commercial relationships.

On scalability: the current architecture cannot, as built, scale to millions of jobs. A single ~54 MB jobs.json at 156,000 jobs implies roughly 350 MB at 1 million and over 1 GB at 3 million — both the storage model and the frontend's whole-file load pattern would need to change. It can, however, comfortably grow into the low hundreds of thousands of jobs with the roadmap in this report, and the existing modular source design makes the eventual move to a datastore-backed, incrementally synced, queryable service a natural evolution rather than a rewrite.